

METIS

Intro to SQL

Joins

Querying and Filtering



```
SELECT name, quantity, vendor_unit_price
FROM inventory
WHERE quantity >= 10
AND vendor_unit_price > 12;
```

name	quantity	vendor_unit_price
giraffe-print bag	18	24.9900
elephant tie	15	13.1900

Querying and Conditioning



VALUE	Vendor Unit Price
High	More than \$20
Medium	Between \$10 and \$20
Low	\$10 or Less

```
SELECT name, vendor_unit_price,  
       CASE  
         WHEN vendor_unit_price IS NULL THEN NULL  
         WHEN vendor_unit_price > 20 THEN "high"  
         WHEN vendor_unit_price > 10 THEN "medium"  
         ELSE "low"  
       END AS value_class  
FROM inventory;
```

Selecting Top Values, Aliasing



```
SELECT TOP 5 id, name, quantity, vendor_unit_price,  
             quantity*vendor_unit_price AS total_inv_value  
FROM inventory  
ORDER BY total_inv_value DESC;
```

id	name	quantity	vendor_unit_price	total_inv_value
2	giraffe-print bag	18	24.9900	449.8200
3	elephant tie	15	13.1900	197.8500
9	bedding set, tiger icons	5	31.9900	159.9500
10	wooly mammoth curtains	4	29.9900	119.9600
4	zebra-striped pants	7	16.8800	118.1600

Aggregating, Filtering, Wildcards



Question:

What is the least expensive item made from each material that starts with the letter 'c'?

```
SELECT material,  
       MIN(sales_price) AS min_price  
FROM item_details  
GROUP BY material  
HAVING material LIKE 'c%';
```

material	min_price
canvas	49.9900
cotton blend	9.9900



JOINS I: CROSS JOIN

METIS

Joining Data Tables



Question:

How much markup does the company charge for each item in stock?

id	name	quantity	vendor_unit_price	last_shipment	reorder
1	tiger t-shirt	10	4.2500	2018-01-22	1
2	giraffe-print bag	18	24.9900	2018-02-26	0
3	elephant tie	15	13.1900	2018-02-26	0
4	zebra-striped pants	7	16.8800		
5	peacock feather hat	2	NULL		
6	leopard-print scarf	NULL	8.5500		
7	walrus-shaped pillow	5	12.2500		
8	gazelle lamp	3	38.8500		
9	bedding set, tiger icons	5	31.9900		
10	wooly mammoth curtains	4	29.9900		

item_id	name	department	material	sales_price
1	tiger t-shirt	clothing	cotton blend	9.9900
2	giraffe-print bag	accessories	canvas	49.9900
3	elephant tie	accessories	silk	35.4900
4	zebra-striped pants	clothing	silk	30.9900
5	peacock feather hat	accessories	felt	34.9900
6	leopard-print scarf	accessories	silk	14.4900
8	gazelle lamp	home goods	metal	79.9900
9	bedding set, tiger icons	home goods	cotton blend	69.9900
11	aardvark earrings	accessories	metal	9.9900

Cross Join



Cartesian Product - product of two sets; yields all possible ordered combinations (a, b)

SET A

2

5

6

**CARTESIAN
PRODUCT**

(2, 'apple') (2, 'banana')

(5, 'apple') (5, 'banana')

(6, 'apple') (6, 'banana')

SET B

'apple'

'banana'

Cross Join Example



Books Table

book_id	book_title	book_genre
1	A Wrinkle in Time	science fiction
2	Murder on the Orient Express	mystery
3	Jurassic Park	science fiction
4	Pride and Prejudice	romance

Directors Table

director_id	director_name	director_specialty	preferred_book_id
1	Alfred Hitchcock	mystery	2
2	Michael Bay	action	3
3	George Lucas	science fiction	1

Cross Join Example



Cartesian Product: Books x Directors

```
SELECT * FROM books, directors;
```

book_id	book_title	book_genre	director_id	director_name	director_specialty	preferred_book_id
1	A Wrinkle in Time	science fiction	1	Alfred Hitchcock	mystery	2
2	Murder on the Orient Express	mystery	1	Alfred Hitchcock	mystery	2
3	Jurassic Park	science fiction	1	Alfred Hitchcock	mystery	2
4	Pride and Prejudice	romance	1	Alfred Hitchcock	mystery	2
1	A Wrinkle in Time	science fiction	2	Michael Bay	action	3
2	Murder on the Orient Express	mystery	2	Michael Bay	action	3
3	Jurassic Park	science fiction	2	Michael Bay	action	3
4	Pride and Prejudice	romance	2	Michael Bay	action	3
1	A Wrinkle in Time	science fiction	3	George Lucas	science fiction	1
2	Murder on the Orient Express	mystery	3	George Lucas	science fiction	1
3	Jurassic Park	science fiction	3	George Lucas	science fiction	1
4	Pride and Prejudice	romance	3	George Lucas	science fiction	1

Note: Cartesian product is also called a CROSS JOIN.

Cross Join Filtered



Question:

What are the titles of the books the directors would like to make into a movie?

```
SELECT * FROM books, directors
WHERE book_id = preferred_book_id;
```

Implicit
Inner Join

book_id	book_title	book_genre	director_id	director_name	director_specialty	preferred_book_id
2	Murder on the Orient Express	mystery	1	Alfred Hitchcock	mystery	2
3	Jurassic Park	science fiction	2	Michael Bay	action	3
1	A Wrinkle in Time	science fiction	3	George Lucas	science fiction	1

Note: Book IDs now match, but not sorted.

Cartesian product created first then rows eliminated with WHERE clause.

JOINS II: INNER, OUTER, LEFT, RIGHT

METIS

Inner Join (Explicit)



Question:

What are the titles of the books the directors would like to make into a movie?

Columns to be
select from
either table

Inner join is the
default type of
join

```
SELECT * FROM books      ← First table
JOIN directors           ← Second table
ON book_id = preferred_book_id;
```

**ON clause provides filtering;
functions similar to WHERE**

book_id	book_title	book_genre	director_id	director_name	director_specialty	preferred_book_id
2	Murder on the Orient Express	mystery	1	Alfred Hitchcock	mystery	2
3	Jurassic Park	science fiction	2	Michael Bay	action	3
1	A Wrinkle in Time	science fiction	3	George Lucas	science fiction	1

Inner Join



ON book_id = preferred_book_id

book_id	book_title	book_genre	director_id	director_name	director_specialty	preferred_book_id
1	A Wrinkle in Time	science fiction	1	Alfred Hitchcock	mystery	2
2 ★	Murder on the Orient Express	mystery	1	Alfred Hitchcock	mystery	2
3	Jurassic Park	science fiction	1	Alfred Hitchcock	mystery	2
4	Pride and Prejudice	romance	1	Alfred Hitchcock	mystery	2
1	A Wrinkle in Time	science fiction	2	Michael Bay	action	3
2	Murder on the Orient Express	mystery	2	Michael Bay	action	3
3 ★	Jurassic Park	science fiction	2	Michael Bay	action	3
4	Pride and Prejudice	romance	2	Michael Bay	action	3
1 ★	A Wrinkle in Time	science fiction	3	George Lucas	science fiction	1
2	Murder on the Orient Express	mystery	3	George Lucas	science fiction	1
3	Jurassic Park	science fiction	3	George Lucas	science fiction	1
4	Pride and Prejudice	romance	3	George Lucas	science fiction	1

Inner Join



Question:

How much markup does the animal items company charge for each item in stock?

id	name	quantity	vendor_unit_price	last_shipment	reorder
1	tiger t-shirt	10	4.2500	2018-01-22	1
2	giraffe-print bag	18	24.9900	2018-02-26	0
3	elephant tie	15	13.1900	2018-02-26	0
4	zebra-striped pants	7	16.8800	2018-02-26	0
5	peacock feather hat	2	NULL	2018-02-26	0
6	leopard-print scarf	NULL	8.5500	2018-02-26	0
7	walrus-shaped pillow	5	12.2500	2018-02-26	0
8	gazelle lamp	3	38.8500	2018-02-26	0
9	bedding set, tiger icons	5	31.9900	2018-02-26	0
10	wooly mammoth curtains	4	29.9900	2018-02-26	0

Join these tables on ID
with explicit inner join

item_id	name	department	material	sales_price
1	tiger t-shirt	clothing	cotton blend	9.9900
2	giraffe-print bag	accessories	canvas	49.9900
3	elephant tie	accessories	silk	35.4900
4	zebra-striped pants	clothing	silk	30.9900
5	peacock feather hat	accessories	felt	34.9900
6	leopard-print scarf	accessories	silk	14.4900
8	gazelle lamp	home goods	metal	79.9900
9	bedding set, tiger icons	home goods	cotton blend	69.9900
11	aardvark earrings	accessories	metal	9.9900

Inner Join



Explicitly tell SQL which table each column comes from: `table.column`

```
SELECT inventory.id, inventory.name,  
       inventory.vendor_unit_price, item_details.sales_price  
FROM inventory  
INNER JOIN item_details  
ON inventory.id = item_details.item_id;
```

Explicitly stating
INNER JOIN also
acceptable

id	name	vendor_unit_price	sales_price
1	tiger t-shirt	4.2500	9.9900
2	giraffe-print bag	24.9900	49.9900
3	elephant tie	13.1900	35.4900
4	zebra-striped pants	16.8800	30.9900
5	peacock feather hat	NULL	34.9900
6	leopard-print scarf	8.5500	14.4900
8	gazelle lamp	38.8500	79.9900
9	bedding set, tiger icons	31.9900	69.9900

Inner Join



```
SELECT inventory.id, inventory.name,  
       inventory.vendor_unit_price, item_details.sales_price,  
       item_details.sales_price-inventory.vendor_unit_price AS markup  
FROM inventory  
INNER JOIN item_details  
ON inventory.id = item_details.item_id ORDER BY markup DESC;
```

id	name	vendor_unit_price	sales_price	markup
8	gazelle lamp	38.8500	79.9900	41.1400
9	bedding set, tiger icons	31.9900	69.9900	38.0000
2	giraffe-print bag	24.9900	49.9900	25.0000
3	elephant tie	13.1900	35.4900	22.3000
4	zebra-striped pants	16.8800	30.9900	14.1100
6	leopard-print scarf	8.5500	14.4900	5.9400
1	tiger t-shirt	4.2500	9.9900	5.7400
5	peacock feather hat	NULL	34.9900	NULL

Joining Tables



Inventory Table

id	name	vendor_unit_price
1	tiger t-shirt	4.2500
2	giraffe-print bag	24.9900
3	elephant tie	13.1900
4	zebra-striped pants	16.8800
5	peacock feather hat	NULL
6	leopard-print scarf	8.5500
7	walrus-shaped pillow	12.2500
8	gazelle lamp	38.8500
9	bedding set, tiger icons	31.9900
10	wooly mammoth curtains	29.9900

No information about item 11

Item Details Table

item_id	name	sales_price
1	tiger t-shirt	9.9900
2	giraffe-print bag	49.9900
3	elephant tie	35.4900
4	zebra-striped pants	30.9900
5	peacock feather hat	34.9900
6	leopard-print scarf	14.4900
8	gazelle lamp	79.9900
9	bedding set, tiger icons	69.9900
11	aardvark earrings	9.9900

No information about items 7 and 10

**INNER JOIN must have information in both tables.
Drops NULL rows from Cartesian product when filtering with ON.**

Outer Join



Unlike INNER JOIN, OUTER JOIN keeps rows where NULL appears in ON clause

Three types of OUTER JOINS to specify which NULL values to keep

1. **FULL OUTER JOIN** – keep everything
2. **LEFT OUTER JOIN** – keep all rows from “left” table, regardless if they appear in “right”
3. **RIGHT OUTER JOIN** – keep all rows from “right” table, regardless if they appear in “left”

FULL OUTER JOIN



```
SELECT inventory.id, inventory.name,  
       inventory.vendor_unit_price,  
       item_details.item_id, item_details.name,  
       item_details.sales_price  
FROM inventory  
FULL OUTER JOIN item_details  
ON inventory.id = item_details.item_id;
```

Only difference now: Specify FULL OUTER JOIN

FULL OUTER JOIN



NULL values appear in both columns from the ON statement

id	name	vendor_unit_price	item_id	name	sales_price
1	tiger t-shirt	4.2500	1	tiger t-shirt	9.9900
2	giraffe-print bag	24.9900	2	giraffe-print bag	49.9900
3	elephant tie	13.1900	3	elephant tie	35.4900
4	zebra-striped pants	16.8800	4	zebra-striped pants	30.9900
5	peacock feather hat	NULL	5	peacock feather hat	34.9900
6	leopard-print scarf	8.5500	6	leopard-print scarf	14.4900
7	walrus-shaped pillow	12.2500	NULL	NULL	NULL
8	gazelle lamp	38.8500	8	gazelle lamp	79.9900
9	bedding set, tiger icons	31.9900	9	bedding set, tiger icons	69.9900
10	wooly mammoth curtains	29.9900	NULL	NULL	NULL
NULL	NULL	NULL	11	aardvark earrings	9.9900

LEFT OUTER JOIN



```
SELECT inventory.id, inventory.name,  
       inventory.vendor_unit_price,  
       item_details.item_id, item_details.name,  
       item_details.sales_price  
FROM inventory  
LEFT OUTER JOIN item_details  
ON inventory.id = item_details.item_id;
```

Left OUTER JOIN will require values for the “left” table (inventory)

LEFT OUTER JOIN



All data from “left” table is retained; NULL values allowed in the id column of the “right” table

id	name	vendor_unit_price	item_id	name	sales_price
1	tiger t-shirt	4.2500	1	tiger t-shirt	9.9900
2	giraffe-print bag	24.9900	2	giraffe-print bag	49.9900
3	elephant tie	13.1900	3	elephant tie	35.4900
4	zebra-striped pants	16.8800	4	zebra-striped pants	30.9900
5	peacock feather hat	NULL	5	peacock feather hat	34.9900
6	leopard-print scarf	8.5500	6	leopard-print scarf	14.4900
7	walrus-shaped pillow	12.2500	NULL	NULL	NULL
8	gazelle lamp	38.8500	8	gazelle lamp	79.9900
9	bedding set, tiger icons	31.9900	9	bedding set, tiger icons	69.9900
10	wooly mammoth curtains	29.9900	NULL	NULL	NULL

Joining Not on Primary Keys



Question:

What possible book + director combinations align in genre-specialty?

Books Table

book_id	book_title	book_genre
1	A Wrinkle in Time	science fiction
2	Murder on the Orient Express	mystery
3	Jurassic Park	science fiction
4	Pride and Prejudice	romance

Directors Table

director_id	director_name	director_specialty	preferred_book_id
1	Alfred Hitchcock	mystery	2
2	Michael Bay	action	3
3	George Lucas	science fiction	1

Joining Not on Primary Keys



Question:

What possible book + director combinations align in genre-specialty?

```
SELECT books.book_title, directors.director_name,  
       directors.director_specialty  
FROM books JOIN directors  
ON books.book_genre = directors.director_specialty  
ORDER BY directors.director_name;
```

book_title	director_name	director_specialty
Murder on the Orient Express	Alfred Hitchcock	mystery
A Wrinkle in Time	George Lucas	science fiction
Jurassic Park	George Lucas	science fiction

Self-Join



Add one additional column to the Directors Table called “fav_director_id”

director_id	director_name	director_specialty	preferred_book_id	fav_director_id
1	Alfred Hitchcock	mystery	2	3
2	Michael Bay	action	3	3
3	George Lucas	science fiction	1	1

Question:

Which other director does each director in our table admire?

Self-Join



Question:

Which other director does each director in our table admire?

```
SELECT directors.director_name,  
       favorite.director_name AS favorite_director  
FROM directors  
JOIN directors AS favorite  
ON directors.fav_director_id = favorite.director_id;
```


director_name	favorite_director
Alfred Hitchcock	George Lucas
Michael Bay	George Lucas
George Lucas	Alfred Hitchcock

**Must alias table to
be able to perform
self-join?**


Types of Joins




SQL JOINS




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key  
WHERE B.Key IS NULL
```




```
SELECT <select_list>  
FROM TableA A  
INNER JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL
```



```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL  
OR B.Key IS NULL
```

**What's the difference between a JOIN
and an INNER JOIN?**

**How much information does a FULL
OUTER JOIN provide?**



QUESTIONS?
